
Cloud AI Model Services Statement Of Work

Aster Cognitive Cloud conversational assistant, document intelligence, retrieval, and model operations services

Synthetic legal-style test document for agent evaluation. Not legal advice. No real parties.

Project Description

Aster Cognitive Cloud will configure a private conversational assistant for enterprise knowledge retrieval, policy question answering, document triage, meeting summary generation, and compliance evidence extraction. The assistant will connect to SharePoint, Microsoft Teams exports, ServiceNow knowledge articles, Jira issues, Confluence spaces, Azure AI Search, and a customer data lake.

Data And Content

- Input data may include policies, procedures, contracts, supplier questionnaires, meeting notes, service tickets, architecture diagrams, audit findings, and selected email exports.
- Personal data may appear in documents, including employee names, email addresses, phone numbers, job titles, issue narratives, approver names, and access logs.
- Highly confidential data, secrets, credentials, private keys, health records, and payment card data must be excluded from indexing unless separately approved.

Model And Data-Use Commitments

- Customer content shall not be used to train, fine-tune, improve, benchmark, or evaluate shared provider models unless a separate opt-in is executed.
- Provider may use customer prompts and outputs only for tenant operations, abuse monitoring, troubleshooting, and quality review approved by customer.
- Provider shall support prompt and response logging controls, retention limits, redaction, and export of model interaction records.

Responsible AI Controls

- Provider shall document intended use, prohibited use, evaluation plan, human review boundary, fallback behavior, hallucination mitigation, and bias-sensitive data handling.
- Provider shall implement retrieval citations, confidence notices, no automatic approval for high-impact decisions, and escalation for legal, HR, finance, safety, or regulated advice.
- Provider shall run scenario-based tests for prompt injection, data exfiltration attempts, unsupported certainty, privacy leakage, and unsafe action recommendations.

Security Architecture

The service will use customer-managed identity provider SSO, tenant-scoped retrieval indexes, private network connectivity where available, key rotation, encrypted storage, API audit logs, and administrative access reviews.

Integrations

- Microsoft Entra ID for authentication and group mapping.
- SharePoint and Teams exports for document retrieval.
- ServiceNow and Jira for ticket knowledge retrieval.
- Azure AI Search or approved equivalent for vector and keyword retrieval.
- SIEM export for audit logs and security monitoring.

Operational Requirements

- Provider shall provide uptime target of 99.5 percent for the assistant API after production acceptance.
- Provider shall support incident response, safety rollback, model version rollback, blocked prompt categories, and emergency kill switch.

Open Preconditions

- Independent bias and robustness test results are not attached.
- Final RAI assessment is pending customer approval.
- Production data retention period is still subject to privacy approval.
- No go-live may occur until human approval, security signoff, and data owner approval are recorded.

Reference Basis For Synthetic Fixture

This fixture is original synthetic text. It uses public legal, regulatory, security, cloud, export-control, and AI governance sources only as structural reference points. It is not legal advice and is not a copy of any source template.

Reference | GOV.UK Model Services Contract | Applied for: complex services structure, schedules, service levels, exit management, security, schedule patterns | <https://www.gov.uk/government/collections/model-services-contract>

Reference | ICO controller and processor contract guidance | Applied for: processor instructions, confidentiality, subprocessors, assistance, deletion, audit, breach support | <https://ico.org.uk/for-organisations/uk-gdpr-guidance-and-resources/accountability-and-governance/contracts-and-liabilities-between-controllers-and-processors-multi/>

Reference | European Commission controller-processor SCCs and transfer SCC overview | Applied for: SCC-style annexes, processing description, technical measures, transfer safeguards | https://commission.europa.eu/publications/standard-contractual-clauses-controller-and-processors-eueea_en

Reference | BIS Export Administration Regulations and advanced computing end-use controls | Applied for: classification, license analysis, end-use, end-user, restricted-party, re-export, transfer controls | <https://www.bis.gov/entity-list>

Reference | NIST SP 800-161 Rev. 1 Cybersecurity Supply Chain Risk Management | Applied for: supplier assurance, provenance, counterfeit/tamper controls, ongoing monitoring, acquisition evidence | <https://csrc.nist.gov/pubs/sp/800/161/r1/upd1/final>

Reference | NIST SP 800-53 Rev. 5 security and privacy control families | Applied for: access control, audit, incident response, contingency planning, acquisition, privacy, supply-chain controls | <https://csrc.nist.gov/pubs/sp/800/53/r5/upd1/final>

Reference | NIST AI Risk Management Framework | Applied for: AI governance, human oversight, testing, documentation, transparency, monitoring, risk management | <https://www.nist.gov/itl/ai-risk-management-framework>

Reference | Cloud Security Alliance Cloud Controls Matrix | Applied for: cloud shared responsibility, identity, logging, data security, resilience, vendor assurance domains | <https://cloudsecurityalliance.org/research/cloud-controls-matrix>

Schedule A - Commercial And Governance Terms

This schedule converts the business request into a production-ready contract package for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction. The parties acknowledge the primary regulated focus as responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring.

1. Definitions And Interpretation

For Cloud AI Model Services, Aster Cognitive Cloud shall define controlled terms, service boundary, covered affiliates, data classes, evidence artifacts, and approval gates. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for definitions and interpretation.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 01 | Area: Definitions And Interpretation | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

2. Order Of Precedence

For Cloud AI Model Services, Aster Cognitive Cloud shall resolve conflicts between the order form, schedules, security annex, DPA, support policies, and statements of work. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for order of precedence.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 02 | Area: Order Of Precedence | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

3. Statement Of Work Controls

For Cloud AI Model Services, Aster Cognitive Cloud shall require every work package to identify scope, owners, data, systems, locations, subcontractors, acceptance tests, and blockers. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for statement of work controls.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 03 | Area: Statement Of Work Controls | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

4. Governance Board

For Cloud AI Model Services, Aster Cognitive Cloud shall establish operational, legal, security, privacy, finance, and compliance reviewers with documented decision rights. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for governance board.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 04 | Area: Governance Board | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

5. Onboarding Gate

For Cloud AI Model Services, Aster Cognitive Cloud shall block production use until risk classification, architecture review, data owner approval, security review, and contract signoff are complete. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for onboarding gate.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 05 | Area: Onboarding Gate | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

6. Service Description

For Cloud AI Model Services, Aster Cognitive Cloud shall state what the provider may do, what it may not do, and which activities need a separate change order. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for service description.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 06 | Area: Service Description | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

7. Customer Responsibilities

For Cloud AI Model Services, Aster Cognitive Cloud shall identify required approvals, access provisioning, data-quality responsibilities, and evidence response timelines. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for customer responsibilities.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 07 | Area: Customer Responsibilities | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

8. Provider Responsibilities

For Cloud AI Model Services, Aster Cognitive Cloud shall require professional care, trained personnel, documented controls, subcontractor oversight, and audit cooperation. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for provider responsibilities.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 08 | Area: Provider Responsibilities | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

9. Regulatory Change

For Cloud AI Model Services, Aster Cognitive Cloud shall require impact assessment, control updates, pricing discussion, and implementation plan when law or regulator guidance changes. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for regulatory change.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 09 | Area: Regulatory Change | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

10. Policies And Standards

For Cloud AI Model Services, Aster Cognitive Cloud shall bind the provider to security, privacy, AI, procurement, continuity, records, and acceptable-use policies supplied by customer. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for policies and standards.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 10 | Area: Policies And Standards | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

11. Personnel And Background Screening

For Cloud AI Model Services, Aster Cognitive Cloud shall require appropriate screening, role-based training, confidentiality obligations, and access removal at exit. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for personnel and background screening.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 11 | Area: Personnel And Background Screening | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

12. Conflicts Of Interest

For Cloud AI Model Services, Aster Cognitive Cloud shall require disclosure of conflicting customers, incentives, related-party arrangements, and independence impairments. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for conflicts of interest.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 12 | Area: Conflicts Of Interest | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

13. Subcontracting

For Cloud AI Model Services, Aster Cognitive Cloud shall require approval for material subcontractors, flow-down clauses, evidence rights, and termination support. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for subcontracting.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 13 | Area: Subcontracting | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

14. Change Control

For Cloud AI Model Services, Aster Cognitive Cloud shall require risk assessment before changing scope, location, subprocessors, integrations, data classes, AI use, or access model. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for change control.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 14 | Area: Change Control | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

15. Fees And Invoicing Controls

For Cloud AI Model Services, Aster Cognitive Cloud shall align billing events to accepted deliverables, usage reports, rate cards, credits, and disputed invoice procedures. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for fees and invoicing controls.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 15 | Area: Fees And Invoicing Controls | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

16. Service Levels

For Cloud AI Model Services, Aster Cognitive Cloud shall define uptime, response, restoration, data delivery, support, reporting, and incident communication measures. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for service levels.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 16 | Area: Service Levels | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

17. Service Credits

For Cloud AI Model Services, Aster Cognitive Cloud shall state credits do not replace corrective actions, regulatory duties, or termination rights for chronic failure. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for service credits.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 17 | Area: Service Credits | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

18. Acceptance Testing

For Cloud AI Model Services, Aster Cognitive Cloud shall require functional, security, privacy, continuity, integration, and reporting tests before production acceptance. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for acceptance testing.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 18 | Area: Acceptance Testing | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

19. Documentation

For Cloud AI Model Services, Aster Cognitive Cloud shall require architecture diagrams, runbooks, integration inventory, data maps, control matrix, and evidence pack maintenance. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for documentation.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 19 | Area: Documentation | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

20. Records Retention

For Cloud AI Model Services, Aster Cognitive Cloud shall preserve approvals, logs, exports, support records, change records, audit evidence, and decision files for defined periods. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for records retention.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 20 | Area: Records Retention | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

21. Audit Rights

For Cloud AI Model Services, Aster Cognitive Cloud shall permit questionnaire, evidence review, independent assurance report, interview, remediation tracking, and targeted audit. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for audit rights.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 21 | Area: Audit Rights | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

22. Incident Management

For Cloud AI Model Services, Aster Cognitive Cloud shall define security, privacy, continuity, physical, export, payment, and data-quality incidents with severity levels. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for incident management.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 22 | Area: Incident Management | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

23. Notification

For Cloud AI Model Services, Aster Cognitive Cloud shall require prompt notices, factual updates, root cause, containment, legal hold, communications approval, and corrective plan. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for notification.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 23 | Area: Notification | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding annual refresh

24. Remediation

For Cloud AI Model Services, Aster Cognitive Cloud shall require tracked corrective actions, owners, deadlines, residual risk acceptance, and executive escalation for overdue blockers. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for remediation.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 24 | Area: Remediation | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding annual refresh

25. Business Continuity

For Cloud AI Model Services, Aster Cognitive Cloud shall require continuity plan, disaster recovery plan, tabletop test, restoration evidence, manual workaround, and dependency mapping. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for business continuity.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 25 | Area: Business Continuity | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

26. Exit Management

For Cloud AI Model Services, Aster Cognitive Cloud shall require transition plan, data export, knowledge transfer, secrets rotation, subcontractor disengagement, and deletion certification. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for exit management.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 26 | Area: Exit Management | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

27. Termination Rights

For Cloud AI Model Services, Aster Cognitive Cloud shall allow termination for uncured breach, security failure, regulatory prohibition, chronic SLA failure, or unresolved approval blocker. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for termination rights.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 27 | Area: Termination Rights | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

28. Liability Allocation

For Cloud AI Model Services, Aster Cognitive Cloud shall separate ordinary commercial liability from confidentiality, data protection, security, fraud, wilful misconduct, and IP claims. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for liability allocation.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 28 | Area: Liability Allocation | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

29. Insurance

For Cloud AI Model Services, Aster Cognitive Cloud shall require technology E&O, cyber, crime, professional liability, transit/warehouse coverage where relevant, and evidence of renewal. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for insurance.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 29 | Area: Insurance | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

30. Confidentiality

For Cloud AI Model Services, Aster Cognitive Cloud shall cover business, technical, pricing, architecture, security, data, model, and investigation information with survival terms. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for confidentiality.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 30 | Area: Confidentiality | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

31. Intellectual Property

For Cloud AI Model Services, Aster Cognitive Cloud shall separate background IP, deliverables, configurations, customer data, trained artifacts, documentation, and feedback. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for intellectual property.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 31 | Area: Intellectual Property | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

32. Open Source And Third-Party Components

For Cloud AI Model Services, Aster Cognitive Cloud shall require inventory, license compliance, vulnerability management, provenance, and replacement rights. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for open source and third-party components.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 32 | Area: Open Source And Third-Party Components | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

33. Publicity Restrictions

For Cloud AI Model Services, Aster Cognitive Cloud shall prohibit public use of customer name, data, logo, case study, benchmark, or regulator discussion without written approval. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for publicity restrictions.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 33 | Area: Publicity Restrictions | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

34. Dispute Escalation

For Cloud AI Model Services, Aster Cognitive Cloud shall require operational escalation, executive escalation, interim performance, evidence preservation, and venue selection. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for dispute escalation.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 34 | Area: Dispute Escalation | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

35. Regulator Cooperation

For Cloud AI Model Services, Aster Cognitive Cloud shall require assistance for supervisory authority inquiries, audits, information requests, breach notifications, and remediation evidence. The obligation applies to private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and to all records concerning retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

- Required evidence: owner, approval date, artifact ID, renewal date, exception status, and reviewer for regulator cooperation.
- Approval blocker: unresolved gaps affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan must remain visible in the risk register until accepted by Customer.

Clause 35 | Area: Regulator Cooperation | Customer: Praxis Knowledge Services | Provider: Aster Cognitive Cloud | Evidence cadence: onboarding and annual refresh

Schedule B - Security, Privacy, Operational, And Supply-Chain Control Matrix

The following controls are written as test evidence for an enterprise compliance agent. Each row should be parsed as a possible obligation, blocker, owner action, or evidence requirement for Cloud AI Model Services.

CTRL-01 Identity And Access

The parties assign identity and access as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering named accounts, MFA, least privilege, joiner-mover-leaver workflow, and quarterly recertification.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-01 | Control: Identity And Access | Evidence: named accounts, MFA, least privilege, joiner-mover-leaver workflow, and quarterly recertification | Blocker if missing: yes for high-risk or regulated processing

CTRL-02 Privileged Access

The parties assign privileged access as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering ticketed elevation, time limits, session recording, command logging, approval evidence, and emergency break-glass review.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-02 | Control: Privileged Access | Evidence: ticketed elevation, time limits, session recording, command logging, approval evidence, and emergency break-glass review | Blocker if missing: yes for high-risk or regulated processing

CTRL-03 Authentication

The parties assign authentication as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering SSO, certificate or key rotation, conditional access, device posture, and credential vaulting.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-03 | Control: Authentication | Evidence: SSO, certificate or key rotation, conditional access, device posture, and credential vaulting | Blocker if missing: yes for high-risk or regulated processing

CTRL-04 Authorization

The parties assign authorization as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering role definitions, segregation of duties, service account ownership, API scope limits, and customer revocation rights.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-04 | Control: Authorization | Evidence: role definitions, segregation of duties, service account ownership, API scope limits, and customer revocation rights | Blocker if missing: yes for high-risk or regulated processing

CTRL-05 Encryption

The parties assign encryption as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering TLS for transit, encryption at rest, key ownership, key rotation, and secure backup encryption.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-05 | Control: Encryption | Evidence: TLS for transit, encryption at rest, key ownership, key rotation, and secure backup encryption | Blocker if missing: yes for high-risk or regulated processing

CTRL-06 Logging

The parties assign logging as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering admin activity, API calls, data exports, model calls, configuration changes, and failed access attempts.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-06 | Control: Logging | Evidence: admin activity, API calls, data exports, model calls, configuration changes, and failed access attempts | Blocker if missing: yes for high-risk or regulated processing

CTRL-07 Monitoring

The parties assign monitoring as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering alert thresholds, SIEM export, anomaly detection, health checks, and operational dashboards.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-07 | Control: Monitoring | Evidence: alert thresholds, SIEM export, anomaly detection, health checks, and operational dashboards | Blocker if missing: yes for high-risk or regulated processing

CTRL-08 Vulnerability Management

The parties assign vulnerability management as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering scanning, severity rules, patch timelines, exception approval, and retest evidence.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-08 | Control: Vulnerability Management | Evidence: scanning, severity rules, patch timelines, exception approval, and retest evidence | Blocker if missing: yes for high-risk or regulated processing

CTRL-09 Secure Development

The parties assign secure development as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering secure SDLC, code review, dependency scanning, branch protection, secrets scanning, and release approval.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-09 | Control: Secure Development | Evidence: secure SDLC, code review, dependency scanning, branch protection, secrets scanning, release approval | Blocker if missing: yes for high-risk or regulated processing

CTRL-10 Penetration Testing

The parties assign penetration testing as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering annual testing, scope approval, remediation evidence, executive summary, and retest letter.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-10 | Control: Penetration Testing | Evidence: annual testing, scope approval, remediation evidence, executive summary, and retest letter | Blocker if missing: yes for high-risk or regulated processing

CTRL-11 Network Security

The parties assign network security as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering segmentation, private connectivity where available, firewall rules, allowlists, and egress control.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-11 | Control: Network Security | Evidence: segmentation, private connectivity where available, firewall rules, allowlists, and egress control | Blocker if missing: yes for high-risk or regulated processing

CTRL-12 Endpoint Security

The parties assign endpoint security as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering managed devices, EDR, disk encryption, patching, removable media restrictions, and lost-device handling.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-12 | Control: Endpoint Security | Evidence: managed devices, EDR, disk encryption, patching, removable media restrictions, and lost device handling | Blocker if missing: yes for high-risk or regulated processing

CTRL-13 Cloud Configuration

The parties assign cloud configuration as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering baseline hardening, configuration drift, infrastructure-as-code review, and tenant isolation verification.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-13 | Control: Cloud Configuration | Evidence: baseline hardening, configuration drift, infrastructure-as-code review, and tenant isolation verification | Blocker if missing: yes for high-risk or regulated processing

CTRL-14 Data Classification

The parties assign data classification as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering classification labels, handling rules, field mapping, retention rules, and data owner signoff.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-14 | Control: Data Classification | Evidence: classification labels, handling rules, field mapping, retention rules, and data ownership signoff | Blocker if missing: yes for high-risk or regulated processing

CTRL-15 Data Minimization

The parties assign data minimization as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering purpose limitation, field reduction, sampling controls, masking, and prohibited-data scanning.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-15 | Control: Data Minimization | Evidence: purpose limitation, field reduction, sampling controls, masking, and prohibited-data scanning | Blocker if missing: yes for high-risk or regulated processing

CTRL-16 Data Retention

The parties assign data retention as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering retention schedule, legal hold, deletion queue, backup retirement, and deletion certificate.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-16 | Control: Data Retention | Evidence: retention schedule, legal hold, deletion queue, backup retirement, and deletion certificate | Blocker if missing: yes for high-risk or regulated processing

CTRL-17 Data Transfer

The parties assign data transfer as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering approved locations, transfer safeguards, encryption, access controls, and destination inventory.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-17 | Control: Data Transfer | Evidence: approved locations, transfer safeguards, encryption, access controls, and destination inventory | Blocker if missing: yes for high-risk or regulated processing

CTRL-18 Subprocessor Oversight

The parties assign subprocessor oversight as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering due diligence, flow-down terms, change notice, objection right, and annual review.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-18 | Control: Subprocessor Oversight | Evidence: due diligence, flow-down terms, change notice, objection right, and annual review
Blocker if missing: yes for high-risk or regulated processing

CTRL-19 Incident Response

The parties assign incident response as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering playbooks, severity taxonomy, containment, forensic preservation, and communication approval.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-19 | Control: Incident Response | Evidence: playbooks, severity taxonomy, containment, forensic preservation, and communication approval | Blocker if missing: yes for high-risk or regulated processing

CTRL-20 Continuity

The parties assign continuity as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering RTO/RPO, tabletop test, alternate processing, dependency mapping, and restoration evidence.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-20 | Control: Continuity | Evidence: RTO/RPO, tabletop test, alternate processing, dependency mapping, and restoration evidence
Blocker if missing: yes for high-risk or regulated processing

CTRL-21 Exit

The parties assign exit as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering export format, assistance period, knowledge transfer, data deletion, and transition governance.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-21 | Control: Exit | Evidence: export format, assistance period, knowledge transfer, data deletion, and transition governance | Blocker if missing: yes for high-risk or regulated processing

CTRL-22 Audit Evidence

The parties assign audit evidence as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering control owner, evidence artifact, refresh cadence, reviewer, deficiency handling, and evidence ID.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-22 | Control: Audit Evidence | Evidence: control owner, evidence artifact, refresh cadence, reviewer, deficiency handling, and evidence ID | Blocker if missing: yes for high-risk or regulated processing

CTRL-23 Responsible AI

The parties assign responsible ai as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering intended use, prohibited use, human approval, testing, monitoring, and unsupported certainty controls.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-23 | Control: Responsible AI | Evidence: intended use, prohibited use, human approval, testing, monitoring, and unsupported certainty controls | Blocker if missing: yes for high-risk or regulated processing

CTRL-24 Model Data Use

The parties assign model data use as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering training opt-in, fine-tuning prohibition, prompt/output retention, service improvement controls, and deletion.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-24 | Control: Model Data Use | Evidence: training opt-in, fine-tuning prohibition, prompt/output retention, service improvement controls, and deletion | Blocker if missing: yes for high-risk or regulated processing

CTRL-25 Integration Governance

The parties assign integration governance as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering system owner, data owner, API scopes, field mapping, failure handling, and rollback.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-25 | Control: Integration Governance | Evidence: system owner, data owner, API scopes, field mapping, failure handling, and rollback | Blocker if missing: yes for high-risk or regulated processing

CTRL-26 Financial Controls

The parties assign financial controls as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering approval authority, segregation of duties, payment-change controls, reconciliation, and exception reporting.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-26 | Control: Financial Controls | Evidence: approval authority, segregation of duties, payment-change controls, reconciliation, exception reporting | Blocker if missing: yes for high-risk or regulated processing

CTRL-27 Supplier Risk

The parties assign supplier risk as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering criticality, dependency, concentration risk, ownership, subcontractor risk, and remediation tracking.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-27 | Control: Supplier Risk | Evidence: criticality, dependency, concentration risk, ownership, subcontractor risk, and remediation tracking | Blocker if missing: yes for high-risk or regulated processing

CTRL-28 Physical Security

The parties assign physical security as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering facility access, visitor logs, chain of custody, asset tracking, and environmental controls.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-28 | Control: Physical Security | Evidence: facility access, visitor logs, chain of custody, asset tracking, and environmental controls | Blocker if missing: yes for high-risk or regulated processing

CTRL-29 Export Controls

The parties assign export controls as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering classification, license analysis, sanctions screening, end-use, end-user, re-export, and transfer restrictions.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-29 | Control: Export Controls | Evidence: classification, license analysis, sanctions screening, end-use, end-user, re-export, and transfer restrictions | Blocker if missing: yes for high-risk or regulated processing

CTRL-30 Marketing Consent

The parties assign marketing consent as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering lawful basis, consent source, suppression, audience upload approval, data partner review, and deletion.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.

- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

CTRL-30 | Control: Marketing Consent | Evidence: lawful basis, consent source, suppression, audience upload approval, data partner review and deletion | Blocker if missing: yes for high-risk or regulated processing

Schedule C - Domain-Specific Requirements For Cloud AI Model Services

These clauses are intentionally specific to Cloud AI Model Services and should trigger specialized reasoning rather than generic vendor-risk handling.

DOM-01 AI System Card

The parties assign ai system card as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering intended use, users, prohibited use, model family, retrieval scope, limitations, and fallback behavior.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

DOM-01 | Control: AI System Card | Evidence: intended use, users, prohibited use, model family, retrieval scope, limitations, and fallback behavior | Blocker if missing: yes for high-risk or regulated processing

DOM-02 Prompt Injection Testing

The parties assign prompt injection testing as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering malicious instructions, data exfiltration, tool abuse, citation spoofing, and unsafe recommendation tests.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

DOM-02 | Control: Prompt Injection Testing | Evidence: malicious instructions, data exfiltration, tool abuse, citation spoofing, and unsafe recommendation tests | Blocker if missing: yes for high-risk or regulated processing

DOM-03 Retrieval Quality

The parties assign retrieval quality as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering index source, freshness, permission trimming, citation quality, confidence indicators, and stale-content handling.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

DOM-03 | Control: Retrieval Quality | Evidence: index source, freshness, permission trimming, citation quality, confidence indicators, stale-content handling | Blocker if missing: yes for high-risk or regulated processing

DOM-04 Human Oversight

The parties assign human oversight as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering approval boundary, escalation queues, reviewer role, override record, and no autonomous high-impact decision.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

DOM-04 | Control: Human Oversight | Evidence: approval boundary, escalation queues, reviewer role, override record, and no autonomous impact decision | Blocker if missing: yes for high-risk or regulated processing

DOM-05 Model Operations

The parties assign model operations as a control family for responsible AI, privacy, prompt injection, retrieval quality, auditability, human oversight, security monitoring. Provider must operate the control for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction and retain evidence covering versioning, rollback, gateway logging, rate limits, retention, redaction, and incident response.

- Testing procedure: sample the latest production configuration, one recent change, one access approval, one exception, and one renewal artifact.
- Failure mode: if evidence is stale, incomplete, or inconsistent with contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, the control is treated as not operating for approval purposes.

DOM-05 | Control: Model Operations | Evidence: versioning, rollback, gateway logging, rate limits, retention, redaction, and incident response | Blocker if missing: yes for high-risk or regulated processing

Schedule D - Operational Runbook And Scenario Tests

Provider shall maintain a runbook for normal operations, failed operations, emergency operations, and exit operations involving retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs.

D.1 Scenario - New Production Launch

If a new production launch occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-01 | Trigger: new production launch | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.2 Scenario - Material Subprocessor Change

If a material subprocessor change occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-02 | Trigger: material subprocessor change | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.3 Scenario - Privileged Support Request

If a privileged support request occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-03 | Trigger: privileged support request | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.4 Scenario - Suspected Data Leakage

If a suspected data leakage occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-04 | Trigger: suspected data leakage | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.5 Scenario - Regional Outage

If a regional outage occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-05 | Trigger: regional outage | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.6 Scenario - Billing Dispute

If a billing dispute occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-06 | Trigger: billing dispute | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.7 Scenario - Regulatory Inquiry

If a regulatory inquiry occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-07 | Trigger: regulatory inquiry | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.8 Scenario - Failed Access Review

If a failed access review occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-08 | Trigger: failed access review | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.9 Scenario - Expired Assurance Report

If a expired assurance report occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-09 | Trigger: expired assurance report | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.10 Scenario - Unapproved Integration

If a unapproved integration occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-10 | Trigger: unapproved integration | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.11 Scenario - Customer Exit

If a customer exit occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-11 | Trigger: customer exit | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.12 Scenario - Security Incident

If a security incident occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-12 | Trigger: security incident | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.13 Scenario - Privacy Incident

If a privacy incident occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-13 | Trigger: privacy incident | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.14 Scenario - Change Rollback

If a change rollback occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-14 | Trigger: change rollback | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.15 Scenario - Evidence Challenge

If a evidence challenge occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-15 | Trigger: evidence challenge | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.16 Scenario - Policy Exception

If a policy exception occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-16 | Trigger: policy exception | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.17 Scenario - Critical Vulnerability

If a critical vulnerability occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-17 | Trigger: critical vulnerability | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.18 Scenario - Unsupported Component

If a unsupported component occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-18 | Trigger: unsupported component | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.19 Scenario - Operator Departure

If a operator departure occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-19 | Trigger: operator departure | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.20 Scenario - Service Level Breach

If a service-level breach occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-20 | Trigger: service-level breach | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.21 Scenario - Data Deletion Request

If a data deletion request occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-21 | Trigger: data deletion request | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.22 Scenario - Audit Sample Failure

If a audit sample failure occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-22 | Trigger: audit sample failure | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.23 Scenario - Dependency Concentration Issue

If a dependency concentration issue occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-23 | Trigger: dependency concentration issue | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.24 Scenario - Cross Border Transfer Review

If a cross-border transfer review occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-24 | Trigger: cross-border transfer review | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

D.25 Scenario - Third Party Destination Review

If a third-party destination review occurs, Provider shall open a tracked record, identify the affected service boundary, assess impact to contracts, policies, supplier questionnaires, tickets, architecture notes, meeting notes, approver names, and access logs, assign owner and deadline, and preserve evidence for audit review.

- Required decision output: continue, pause, remediate, escalate, or exit. The decision must reference current evidence and any gap affecting independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan.

SCN-25 | Trigger: third-party destination review | Evidence: ticket, owner, chronology, impact, control response, closure approval | Reviewer: customer control owner

Schedule F - Evidence Acceptance Matrix

EVID-001	Artifact: contract approval Applies to: Cloud AI Model Services Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence
EVID-001A	Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.
EVID-001B	Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.
EVID-002	Artifact: risk assessment Applies to: Cloud AI Model Services Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence
EVID-002A	Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.
EVID-002B	Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.
EVID-003	Artifact: data map Applies to: Cloud AI Model Services Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence
EVID-003A	Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.
EVID-003B	Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.
EVID-004	Artifact: system architecture Applies to: Cloud AI Model Services Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence
EVID-004A	Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.
EVID-004B	Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.
EVID-005	Artifact: integration inventory Applies to: Cloud AI Model Services Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence
EVID-005A	Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.
EVID-005B	Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.
EVID-006	Artifact: security assurance Applies to: Cloud AI Model Services Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-006A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-006B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-007 | Artifact: access approval | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-007A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-007B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-008 | Artifact: privileged access review | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-008A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-008B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-009 | Artifact: change record | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-009A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-009B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-010 | Artifact: incident playbook | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-010A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-010B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-011 | Artifact: continuity test | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-011A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-011B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-012 | Artifact: exit plan | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-012A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-012B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-013 | Artifact: subcontractor review | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-013A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-013B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-014 | Artifact: retention schedule | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-014A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-014B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-015 | Artifact: deletion certificate | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-015A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-015B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-016 | Artifact: service level report | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-016A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-016B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-017 | Artifact: availability report | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-017A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-017B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-018 | Artifact: support ticket sample | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-018A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-018B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-019 | Artifact: audit log export | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-019A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-019B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-020 | Artifact: training record | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-020A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-020B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-021 | Artifact: policy exception | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-021A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-021B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-022 | Artifact: threat model | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-022A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-022B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-023 | Artifact: vulnerability report | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-023A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-023B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-024 | Artifact: penetration test summary | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-024A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-024B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-025 | Artifact: remediation plan | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-025A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-025B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-026 | Artifact: business owner signoff | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-026A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-026B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-027 | Artifact: legal approval | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-027A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-027B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-028 | Artifact: privacy approval | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-028A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-028B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-029 | Artifact: security approval | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-029A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-029B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

EVID-030 | Artifact: finance approval | Applies to: Cloud AI Model Services | Minimum fields: owner, date, source system, reviewer, finding status, renewal cadence

EVID-030A | Acceptance rule: artifact must address retrieval indexes, prompts, responses, model gateways, embeddings, connectors, human review queues, and safety logs and must not contradict the stated scope for private assistant, retrieval, document intelligence, policy question answering, meeting summaries, and compliance evidence extraction.

EVID-030B | Blocking rule: if artifact is absent, expired, unsigned, or limited to a non-production environment, the approval status remains conditional.

Schedule G - Known Gaps And Approval Conditions

As of the synthetic effective date, the following items are intentionally incomplete to test blocker detection: independent robustness results, final RAI assessment, retention approval, data owner approval, model rollback plan. No production approval may be inferred from partial evidence.

G.1 Open Item - Independent Robustness Results

The open item 'independent robustness results' remains a blocking or conditional approval factor until a named owner supplies final evidence, legal and security review are complete, and the residual risk decision is recorded.

- Agent expectation: classify the gap, map the affected domain, state the missing artifact, and recommend the next owner action.

G.2 Open Item - Final Rai Assessment

The open item 'final RAI assessment' remains a blocking or conditional approval factor until a named owner supplies final evidence, legal and security review are complete, and the residual risk decision is recorded.

- Agent expectation: classify the gap, map the affected domain, state the missing artifact, and recommend the next owner action.

G.3 Open Item - Retention Approval

The open item 'retention approval' remains a blocking or conditional approval factor until a named owner supplies final evidence, legal and security review are complete, and the residual risk decision is recorded.

- Agent expectation: classify the gap, map the affected domain, state the missing artifact, and recommend the next owner action.

G.4 Open Item - Data Owner Approval

The open item 'data owner approval' remains a blocking or conditional approval factor until a named owner supplies final evidence, legal and security review are complete, and the residual risk decision is recorded.

- Agent expectation: classify the gap, map the affected domain, state the missing artifact, and recommend the next owner action.

G.5 Open Item - Model Rollback Plan

The open item 'model rollback plan' remains a blocking or conditional approval factor until a named owner supplies final evidence, legal and security review are complete, and the residual risk decision is recorded.

- Agent expectation: classify the gap, map the affected domain, state the missing artifact, and recommend the next owner action.